

Backbone-Store: An API for Automating Local Caching of Web Application Data

Michael Stunes, Wolfe Styke, Fan (Frances) Zhang

May 21, 2012

Abstract

Web applications that require re-fetching data from a server can see an improved load time by storing old data in the client’s browser, requiring only new data to be retrieved from the server. In this paper, we present Backbone-Store: a local caching API that extends Backbone.js and uses the HTML5 localStorage API to offer web developers two different persistent strategies that easily store and reload a web application’s data. To evaluate the API, we analyze three simple applications to measure the performance characteristics of loading documents, medium-sized emails, and small notes from both the local storage on the client’s browser and from the server. As expected, we find that our two implementations of local storage both perform faster than loading the same data from the server, with the one-key-per-collection persistence strategy outperforming the one-key-per-model strategy. This time difference increases for larger quantities of the same type of item, suggesting that local storage becomes more advantageous as the amount of data loaded by an application increases. Furthermore, we examine the time required for an application to save data to the local storage, comparing the relative save-time performance of our two persistence strategies and finding one to be more suitable for small updates to the application’s data and the

other to be more suitable for larger updates.

1 Introduction

HTML5 presents new options for storing a web application's data in the client's browser. Efficient local storage presents two major benefits. The first benefit is that web applications can store information that is loaded and saved frequently to help speed up the time it takes to present the user interface. Users are sensitive to the time it takes a web application to load, with slower applications causing users to navigate away or think less of the site. The second benefit is that the web applications can request less data from a server if it can store information that it previously would have needed to fetch from the server. This allows a server to fulfill more requests and thus reduces costs to maintain servers to meet a web application's requests.

2 Related Work

There are several new client-side data storage options for web applications using HTML5. Web SQL provides a database interface similar to SQLite and supports transactions. However, Web SQL is not supported by all of the major browsers, with Firefox in particular refusing to ever support Web SQL. HTML5 also offers two key-value storage options called `sessionStorage` and `localStorage`, which provide a simple API (`setItem(key, value)`, `getItem(key)`, `removeItem(key)`, and `clear()`) for storing key-value pairs. Session storage stores data only while a user is viewing a website, and automatically clears the data when a user navigates away from the site. Local storage persists data until it is deleted by the application or by the user clearing all of the website's information from the browser. In this paper, we use `localStorage` to implement our local caching strategies because it provides

persistent storage (unlike `sessionStorage`) and is supported by all the major web browsers (unlike Web SQL).

Backbone-store, the project described in this paper, presents a new way for web application programmers to easily save their web application's data on the client side to take advantage of the benefits of local storage (as described in the Introduction). This project is an extension of the pure-Javascript Backbone.js [?] library, which is used by many companies (e.g., Four Square, Walmart Mobile, Pandora, Trello, etc.) to create web applications for browsers and mobile devices. Data stored in Backbone.js applications are generally placed in either a Model object or a Collection object, which is an ordered list of objects of one type of Model. Backbone.js also provides View objects for managing the UI based on the state of the Model or Collection objects. We used Backbone.js as a base due to its aforementioned popularity with many web developers, and because it provides a very good model for developers that is easy to modify with our extensions.

Currently, web programmers wishing to employ local caching using Backbone.js would need to write a custom storage solution (including custom libraries and scripts tailored to the particular application) to save the data from each of their models and/or collections to HTML5's `localStorage` or another local storage option. Backbone-store simplifies this process by providing an API that automates the local caching process such that information from a saved model or collection can be reloaded into the models and collections the next time the web application starts. With a few new methods added to the Model and Collection objects, a web developer can choose when a copy of a model's or collection's data should be stored, and when to load saved data back into the application's models or collections. There is also a choice of how the data is stored for collections of models. All models can be stored as a bundle under one key in the `localStorage` (Strategy 1), or they can be stored individually, with each model under its own key in the `localStorage` (Strategy 2). These two persistence strategies provide flexibility for different use cases where one option may be significantly

better than the other, depending on the number of items and the amount of data being stored. We evaluate the performance of these two strategies in the Results section.

3 The Problem Space

The first main problem is the diversity of data types stored by web applications. Indeed, Web applications vary in both the quantity and size of data item they must request from a server before loading the user interface, causing possible variance in the load and save times of various persistence strategies. For example, some applications might require only a few large items, such as showing a user a menu of large documents they can edit. On the opposite end of both the quantity and size spectra, a note-taking application might request many relatively small notes. We address the varying data characteristics by evaluating the load and save times of each persistent strategy for applications using data of varying quantities and sizes and highlight cases where each strategy prevails.

The second problem is the challenge of transforming application data stored by Backbone.js into a form that readily interfaces with HTML5's localStorage API. Whereas Backbone.js stores models and collections as Javascript objects, the HTML5 localStorage requires objects to be represented as string keys and string or integer values. We resolve this problem by using the native JSON parse and stringify methods that convert objects to strings and back again, allowing our API to easily store the data from the models and collections as strings in localStorage, and convert the stored strings back to their original object representations.

4 Solution: Extending Backbone.js

There are two primary ways to save Backbone.js collections consisting of many models. One option is to turn the list of models into a JSON string and store this information under a single key in localStorage (Strategy 1). The other option is to store each model in the collection under its own key (Strategy 2). The main tradeoff in these options is the time it takes to retrieve a single key from localStorage, and the time it takes to parse a JSON string.

Our library extends Backbone.js to provide these two new persistence storage options. For a web developer to use these options, the developer chooses a default persistence strategy for all of their collections of models. Additionally, the developer can choose a persistence strategy for each individual collection. This allows the developer to easily specify a single persistence strategy that will cover the majority of their needs, and only change the strategy for the minority of collections that will use the non-default option.

To store a model or a collection's models, a new instance method has been added to the Backbone.Model and Backbone.Collection options called **"saveToStore"**. For models, this method simply saves the model's attributes to the localStorage. For collections, this method uses the persistence strategy identified by the collection to save the sorted list of models into the localStorage. To load the data for a model or collection from localStorage, a new method called **"loadFromStore"** has been added to Backbone.Model and Backbone.Collection. For models, this simply loads the data from localStorage into the model; for collections, this method identifies the persistence strategy used by the collection and uses the correct retrieval method to fetch the data from localStorage, parse it back into an array of models, and add them to the collection.

5 Contribution

Our project makes two main contributions. First, we create an API that automates local caching for web applications. This API will allow developers to choose from two persistence strategies for locally cached data:

1. Associate one key per collection
2. Associate one key per model in a collection

Second, we conduct a performance evaluation of load times for data stored locally (using the aforementioned two strategies) compared to data stored remotely for three types of web applications that differ in item size per item (70 bytes, 900 bytes, 3200 bytes) and item quantity (10, 50, 100, 200, 300 items per collection) and highlight conditions where local storage is most advantageous.

We hypothesize that Backbone-Store can be used by web applications to increase the speed required to load a web application’s user interface by providing faster access to items of varying quantities and sizes.

6 Evaluation

To evaluate the benefits of Backbone-store, we built three test applications to understand how our solution handles three different sizes (70 bytes, 900 bytes, 3200 bytes) of items requiring storage. To test large items that might represent a document, we used five paragraphs of Lorem ipsum, totaling 3265 characters per item (3200 bytes). To examine medium size items such as emails, we used one paragraph of Lorem ipsum, totaling 905 characters per item (900 bytes) . To examine smaller items such as personal notes, we used a short string of 70 characters (70 bytes). These character lengths are the length of the JSON string that represents each item. The three applications performed the same task:

1. Load some number of a particular type of item (e.g., doc, email, note) from local storage
2. Load some number of a particular type of item (e.g., doc, email, note) from the server
3. Record the load time.

Note that Step 1 and 2 are separated by a delay so as not to overlap in execution and skew the load time.

We created a simple server using Node.js [?] for our test applications. Our timing application presented a blank page which loaded a particular number of items for each of the three types of models (Documents, Emails, Notes) into collections from both local storage and from server requests. Three request URLs allowed the web application to use Ajax to query the server for the data needed for each type of model. To easily collect a large enough sample size of timing data, we had the server record responses from the web application and print out batches of times for 15 requests to a CSV file. This allowed us to rapidly collect timing data for each scenario.

Measuring Load Time Each of the three web applications records how long it takes the application to load the data for each of the three test applications. For the time it takes to load data from the server, the timer is started just before an Ajax request asks the server for the data. When the data returns from the server, it is loaded into a Backbone.Collection object, and the timer is stopped after this is complete. To record the time to load the same data from local storage, the timer is started when just before the `loadFromStore` method is called by a collection. The data is then fetched from local storage and parsed from a JSON string into an array of objects, which is then added to the collection before finally stopping the timer.

Measuring Save Time Finally, we tested save times for our two local storage strategies, by creating collections with different numbers of items and measuring the time to save each

collection to local storage. Also, we believe that our two strategies will behave differently in different situations (i.e., saving an entire collection for the first time, and saving a collection again after having changed or added only a few models). We test this by measuring the time to resave the collection after adding one new model. We predict that the one-key-per-model local storage strategy will allow faster updates, by making it possible to save only those models that were changed or added since the last save.

7 Results

Below we present results that measure the load times of data from both local storage and the server for three types of applications, listed here, in decreasing order of data size:

1. Documents (3200 bytes)
2. Emails (900 bytes)
3. Notes (70 bytes)

We load the data from each application in varying quantities (10, 50, 100, 200, 300 items per collection) and show that local caching performs faster than data loaded from the server (except for collections of small items, where the gain becomes negligible, most likely because of the overhead involved in retrieving items from the local store).

Loading large items, such as documents, from the local storage (shown in Figure 1A) gives the results that we expect. Load times for both local persistence strategies, and from the server, grow linearly with the number of items in the collection, and thus with the amount of data retrieved. Both local persistence methods are faster than retrieving items from the local storage. The one-key-per-model method is slightly slower than the one-key-per-collection method due to the additional overhead of retrieving each model from the local storage separately.

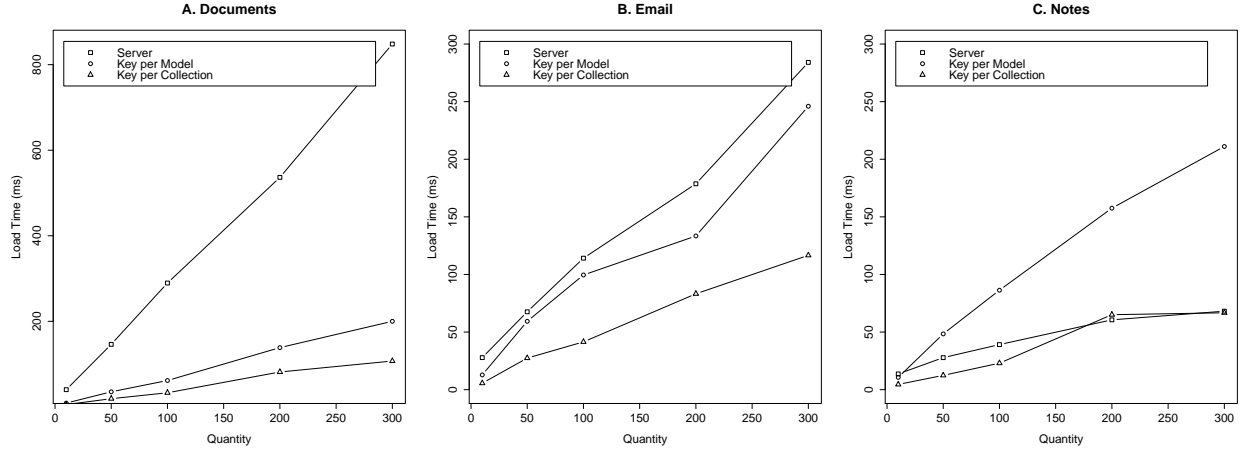


Figure 1: Load time as a function of collection size, for each persistence method and item size.

Repeating the same tests with medium-sized items, such as emails (shown in Figure 1B), shows similar results; however, the difference in speed here is less pronounced. This is due to the constant overhead imposed by accessing the local storage. In the documents example, this overhead is masked by the linear overhead of retrieving large amounts of data from storage. As in the previous example, the one-key-per-model method is still slightly slower, due to additional overhead from retrieving each model from the storage separately.

Finally, repeating these tests with very small items, such as notes (shown in Figure 1C), shows different results: in this case, loading from local storage is negligibly faster, with the difference becoming smaller for larger collections. With very small items in the collection, the additional overhead of the one-key-per-model method becomes large enough to outweigh the load-time benefits of local storage.

It is of note that users typically perceive lag in an interface when an operation takes more than 100 milliseconds to complete. In all of our tests, the one-key-per-collection persistence strategy was able to keep the load time for the collection below or very near 100 ms, meaning that users would perceive the load time as nearly instantaneous. On the other hand, the load time from the server is much higher, so users would perceive noticeable lag while waiting for

the application to load. This demonstrates that our extension provides application users with an improved user experience, due to noticeably faster application load times. Additionally, moving storage from server to client will also reduce the load on the server, allowing it to take more requests from applications, further enhancing its performance and speed.

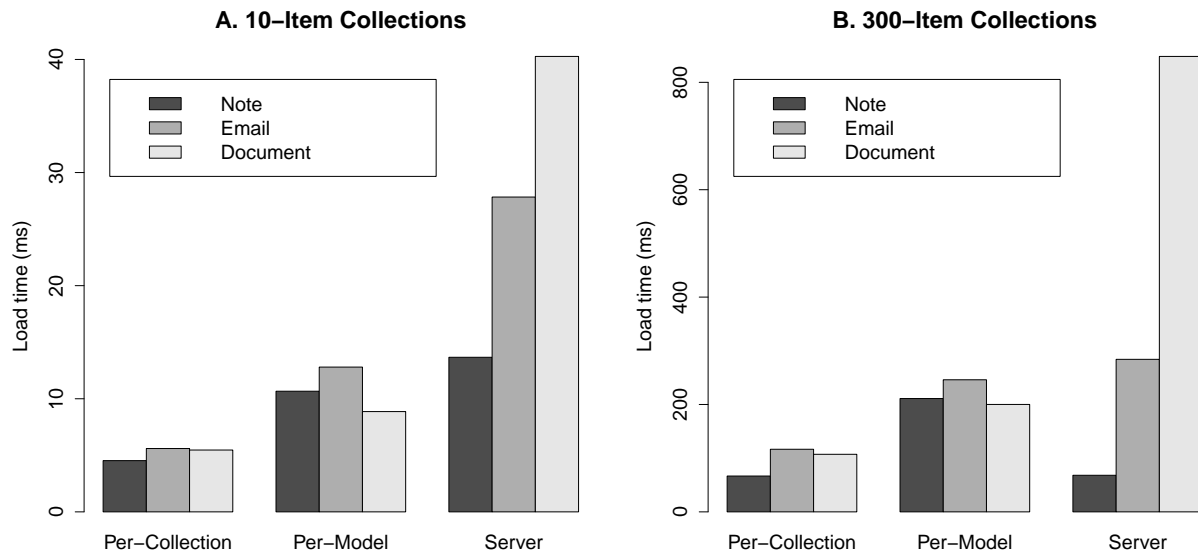


Figure 2: Load time as a function of item size, for each persistence method and two different collection sizes.

Figure 2 shows the effect of the size of individual items on the time to load collections both from storage and from the server. As expected, load times from the server are generally higher. Also, the difference in load time due to the size of the individual items is much more pronounced when loading from the server. Differences in load times due to item size for retrieving collections are much less pronounced when loading from local storage; again, much of the overhead in loading locally is fixed, and caused simply by accessing the local store.

Figure 3 shows the elapsed time for saving a collection of document-size items into local-Storage, while varying the number of items in the collection. For both persistence strategies, the time to save the entire collection grows linearly with the size of the collection as expected;




Figure 3: Save time as a function of collection size, for each persistence method.

furthermore, the one-key-per-model strategy takes longer to save an entire collection, because of the overhead involved in writing many keys to the `localStorage`.

Also, as expected, the time to resave the collection after adding one model is much lower (approx. 1 ms in every case) for the one-key-per-model strategy, as we are able to save only that model to the `localStorage`. For the one-key-per-collection strategy, the time to save one more model is approximately the same as the time to save all of the models, because this strategy must write the entire collection to `localStorage` in either case.

8 Conclusion

This report describes Backbone-Store: our extension for the Backbone.js library, which simplifies the task of persisting a collection of data models to a web browser's local storage, in order to facilitate faster load times and reduce server load. This extension provides an API that automatically saves and loads collections on demand, using one of two different persistence strategies: one-key-per-model and one-key-per-collection. We test our extension using three different demo applications (Documents, Emails, Notes), with different characteristics for the data that are being used by the application. Our tests vary both the number of data items in a single collection and the size of each data item, to simulate a number of different application scenarios. Finally, our test applications show that Backbone-Store is effectively and significantly reduces the load times in many different scenarios, resulting in an improved experience for the application user and a more convenient, simpler programming experience for application developers.